

Chapter 4 Clustering

We will now get into clustering. There are two broad areas of clustering: hard and soft. Hard clustering is when an observation can belong to one and only one cluster. Soft clustering is when an observation receives some type of measure (for example, a probability) of being in a cluster. We will be focusing on hard clustering.

Within hard clustering, algorithms can be flat (for example, kmeans) or hierarchical. Finally, within hierarchical, we have agglomerative and divisive (just like they sound...agglomerative means we will build things up where each observations starts as its own cluster; divisive means we will tear things down which means all observations start in one cluster).

We will start by looking at kmeans, then move into hierarchical clustering.

4.1 Kmeans


We will be using the USArrests data set (this data is already in R). It is an older data set, but still good to illustrate clustering. This data has 50 observations (one observation per state) and contains statistics in arrests per 100,000 residents for assault, murder and rape. Also provided is the percent of the population in that state who live in urban areas.

```
summary(USArrests)
```


##	Murder	Assault	UrbanPop	Rape
##	Min. : 0.800	Min. : 45.0	Min. : 32.00	Min. : 7.30
##	1st Qu.: 4.075	1st Qu.: 109.0	1st Qu.: 54.50	1st Qu.: 15.07
##	Median : 7.250	Median : 159.0	Median : 66.00	Median : 20.10
##	Mean : 7.788	Mean : 170.8	Mean : 65.54	Mean : 21.23
##	3rd Qu.: 11.250	3rd Qu.: 249.0	3rd Qu.: 77.75	3rd Qu.: 26.18
##	Max. : 17.400	Max. : 337.0	Max. : 91.00	Max. : 46.00

```
hist(USArrests$Murder)
```


Histogram of USArrests\$Murder




```
hist(USArrests$Assault)
```



```
hist(USArrests$Rape)
```

Histogram of USArrests\$Rape

```
hist(USArrests$UrbanPop)
```

Histogram of USArrests\$UrbanPop

```
arrest.scal=scale(USArrests)
```

With only 4 variables and good scales, we will just use the scaled data (arrest.scal) for the clustering.


```
clus2=kmeans(arrest.scal,centers=2,nstart = 25)
clus2
```

```
## K-means clustering with 2 clusters of sizes 20, 30
##
## Cluster means:
```

```
##      Murder      Assault      UrbanPop      Rape
## 1  1.004934  1.0138274  0.1975853  0.8469650
## 2 -0.669956 -0.6758849 -0.1317235 -0.5646433
##
## Clustering vector:
##      Alabama      Alaska      Arizona      Arkansas      California
##           1           1           1           2           1
##      Colorado      Connecticut      Delaware      Florida      Georgia
##           1           2           2           1           1
##      Hawaii      Idaho      Illinois      Indiana      Iowa
##           2           2           1           2           2
##      Kansas      Kentucky      Louisiana      Maine      Maryland
##           2           2           1           2           1
##      Massachusetts      Michigan      Minnesota      Mississippi      Missouri
##           2           1           2           1           1
##      Montana      Nebraska      Nevada      New Hampshire      New Jersey
##           2           2           1           2           2
##      New Mexico      New York      North Carolina      North Dakota      Ohio
##           1           1           1           2           2
##      Oklahoma      Oregon      Pennsylvania      Rhode Island      South Carolina
##           2           2           2           2           1
##      South Dakota      Tennessee      Texas      Utah      Vermont
##           2           1           1           2           2
##      Virginia      Washington      West Virginia      Wisconsin      Wyoming
##           2           2           2           2           2
##
## Within cluster sum of squares by cluster:
## [1] 46.74796 56.11445
## (between_SS / total_SS =  47.5 %)
##
## Available components:
```

```
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"
```


```
fviz_cluster(clus2, data = arrest.scal)
```




We do not know if 2 clusters is the best number of clusters, so we will need to explore other options and determine which is the best number of clusters for this data set.

```
set.seed(12964)
```


```
fviz_nbclust(arrest.scal, kmeans, method = "wss", k.max = 9)
```



```
fviz_nbclust(arrest.scal, kmeans, method = "silhouette", k.max = 9)
```

```
set.seed(123)
gap_stat = clusGap(arrest.scal, FUN = kmeans, nstart = 25, K.max = 9, B = 50)
fviz_gap_stat(gap_stat)
```




As you can see, different statistics provide different number of clusters (of course!). As a final thought, let's plot all of these using PCA.

```
k2 <- kmeans(arrest.scal, centers = 2, nstart = 25)
k3 <- kmeans(arrest.scal, centers = 3, nstart = 25)
k4 <- kmeans(arrest.scal, centers = 4, nstart = 25)
```

```
# plots to compare
```

```
p2 <- fviz_cluster(k2, geom = "point", data = arrest.scal) + ggtitle("k = 2")
p3 <- fviz_cluster(k3, geom = "point", data = arrest.scal) + ggtitle("k = 3")
p4 <- fviz_cluster(k4, geom = "point", data = arrest.scal) + ggtitle("k = 4")
```

```
grid.arrange(p2, p3, p4, nrow = 2)
```





Going with 2 clusters, we can now try to profile the data. First, we need to merge the data with the cluster numbers.

```
profile.kmeans=cbind(USArrests,k2$cluster)
all.k=profile.kmeans %>% group_by(k2$cluster) %>%
  summarise(mean.assault=mean(Assault),mean.murder=mean(Murder),mean.rape=mean(I
all.k
```



```
## # A tibble: 2 × 5
```

```
##   `k2$cluster` mean.assault mean.murder mean.rape mean.pop
##           <int>         <dbl>         <dbl>         <dbl>         <dbl>
## 1             1          114.           4.87           15.9           63.6
## 2             2          255.          12.2           29.2           68.4
```

```
NbClust(arrest.scal,method="kmeans",min.nc=2,max.nc = 4)
```



```
## *** : The Hubert index is a graphical method of determining the number of clusters.
##       In the plot of Hubert index, we seek a significant knee that indicates a
##       significant increase of the value of the measure i.e the significant increase in the
##       Hubert index second differences plot.
##
```



```
## *** : The D index is a graphical method of determining the number of clusters.
##           In the plot of D index, we seek a significant knee (the significant
##           second differences plot) that corresponds to a significant increase in
##           the measure.
```

```
##
```

```
## *****
```

```
## * Among all indices:
```

```
## * 13 proposed 2 as the best number of clusters
```

```
## * 9 proposed 3 as the best number of clusters
```

```
## * 2 proposed 4 as the best number of clusters
```

```
##
```

```
##           ***** Conclusion *****
```

```
##
```

```
## * According to the majority rule, the best number of clusters is 2
```

```
##
```

```
##
```

```
## *****
```

```
## $All.index
```

```
##           KL           CH Hartigan           CCC           Scott Marriot           TrCovW           TraceW Friedman
```

```
## 2 5.1512 43.4620 15.0387 -0.1410 81.1595 690539.9 716.8223 102.8624 6.8084
```

```
## 3 2.4059 35.3076 5.6874 -1.0950 124.5851 651900.1 572.6853 78.3233 9.2061
```

```
## 4 0.4884 27.6808 2.5302 -2.6596 152.4673 663556.1 455.6235 69.8686 11.1528
```

```
##           Rubin Cindex           DB Silhouette           Duda Pseudot2           Beale Ratkowsky           Ball
```

```
## 2 1.9055 0.4003 1.0397 0.4085 1.2557 -6.3122 -0.4657 0.4466 51.4312
```

```
## 3 2.5024 0.3595 1.1283 0.3094 0.7004 6.8425 0.9679 0.4417 26.1078
```

```
## 4 2.8053 0.3525 1.1774 0.2590 1.2645 -3.5562 -0.4629 0.3978 17.4671
```

```
##           Ptbiserial           Frey McClain           Dunn Hubert SDindex Dindex           SDbw
```

```
## 2 0.6323 1.7441 0.5470 0.2214 0.0085 1.6841 1.3346 1.1098
```

```
## 3 0.5414 3.4826 1.1571 0.1194 0.0092 1.8410 1.1499 1.4195
```

```
## 4      0.4796 -1.9450  1.6045 0.1235 0.0095  2.1248 1.0671 0.5076
##
## $All.CriticalValues
##   CritValue_Duda CritValue_PseudoT2 Fvalue_Beale
## 2      0.3890      48.6824      1.0000
## 3      0.3508      29.6102      0.4318
## 4      0.2805      43.6156      1.0000
##
## $Best.nc
##           KL      CH Hartigan      CCC      Scott Marriot TrCovW Tracel
## Number_clusters 2.0000 2.000  3.0000 2.000  3.0000      3.00  3.000 3.0000
## Value_Index      5.1512 43.462  9.3513 -0.141 43.4256 50295.68 144.137 16.084
##           Friedman Rubin Cindex      DB Silhouette Duda PseudoT2
## Number_clusters  3.0000 3.0000 4.0000 2.0000      2.0000 2.0000  2.0000
## Value_Index      2.3979 -0.2942 0.3525 1.0397      0.4085 1.2557 -6.3122
##           Beale Ratkowsky Ball PtBiserial Frey McClain Dunn
## Number_clusters 2.0000  2.0000 3.0000      2.0000 3.0000  2.000 2.0000
## Value_Index     -0.4657  0.4466 25.3234      0.6323 3.4826  0.547 0.2214
##           Hubert SDindex Dindex  SDbw
## Number_clusters      0 2.0000      0 4.0000
## Value_Index          0 1.6841      0 0.5076
##
## $Best.partition
##      Alabama      Alaska      Arizona      Arkansas      California
##          1          1          1          2          1
##      Colorado Connecticut Delaware      Florida      Georgia
##          1          2          2          1          1
##      Hawaii      Idaho      Illinois      Indiana      Iowa
##          2          2          1          2          2
##      Kansas      Kentucky Louisiana      Maine      Maryland
##          2          2          1          2          1
```


##	Massachusetts	Michigan	Minnesota	Mississippi	Missouri
##	2	1	2	1	1
##	Montana	Nebraska	Nevada	New Hampshire	New Jersey
##	2	2	1	2	2
##	New Mexico	New York	North Carolina	North Dakota	Ohio
##	1	1	1	2	2
##	Oklahoma	Oregon	Pennsylvania	Rhode Island	South Carolina
##	2	2	2	2	1
##	South Dakota	Tennessee	Texas	Utah	Vermont
##	2	1	1	2	2
##	Virginia	Washington	West Virginia	Wisconsin	Wyoming
##	2	2	2	2	2

4.1.1 K-Means in Python

```

from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from scipy.spatial.distance import cdist
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

arrest_py=r.USArrests
scaler = StandardScaler()
arrest_scal = scaler.fit_transform(arrest_py)

```

```
## C:\PROGRA~3\ANACON~1\lib\site-packages\sklearn\utils\validation.py:623: FutureWarning:
##   if not hasattr(array, "sparse") and array.dtypes.apply(is_sparse).any():
## C:\PROGRA~3\ANACON~1\lib\site-packages\sklearn\utils\validation.py:623: FutureWarning:
##   if not hasattr(array, "sparse") and array.dtypes.apply(is_sparse).any():
```

```
clus_py=KMeans(n_clusters=2, random_state=5687, n_init=25).fit(arrest_scal)
```

```
clus_py.labels_
```

```
## array([1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1,
##        0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0,
##        0, 0, 0, 0, 0, 0])
```

```
inertias = []
```

```
silhouette_coefficients = []
```

```
K=range(2,10)
```

```
for k in K:
```

```
    kmean1 = KMeans(n_clusters=k).fit(arrest_scal)
```

```
    kmean1.fit(arrest_scal)
```


```
    inertias.append(kmean1.inertia_)
```

```
    score = silhouette_score(arrest_scal, kmean1.labels_)
```


```
    silhouette_coefficients.append(score)
```

```
## KMeans(n_clusters=2)
## KMeans(n_clusters=3)
## KMeans(n_clusters=4)
## KMeans(n_clusters=5)
## KMeans(n_clusters=6)
## KMeans(n_clusters=7)
## KMeans()
## KMeans(n_clusters=9)

plt.plot(K, inertias, 'bx-')
plt.xlabel('Values of K')
plt.ylabel('Inertia')
plt.title('The Elbow Method using Inertia')
plt.show()
```



```
plt.plot(K, silhouette_coefficients, 'bx-')
plt.xlabel('Values of K')
plt.ylabel('Silhouette Coefficient')
plt.title('The Silhouette Method')
plt.show()
```



4.2 Hierarchical Clustering


In hierarchical clustering, we need to specify the distance metric we wish to use, as well as the linkage information. There are MANY different distance measures, but the two most common are Euclidean and manhattan. We will use the algorithms in the cluster package in R (provides most flexibility in algorithms). In order to do the basic hierarchical clustering, the algorithm is called agnes (Agglomerative “Nested” Hierarchical Clustering). The code below shows how to do this with complete linkage and using Euclidean distances.

```

dist.assault=dist(arrest.scal,method = "euclidean")
h1.comp.eucl=agnes(dist.assault,method="complete")
pltree(h1.comp.eucl, cex = 0.6, hang = -1, main = "Dendrogram of agnes")

```

Dendrogram of agnes



dist.assault
agnes (*, "complete")

If you want to try a number of different linkage possibilities:

```
m <- c( "average", "single", "complete", "ward")
names(m) <- c( "average", "single", "complete", "ward")
```

```
# function to compute coefficient
```

```
ac <- function(x) {
  agnes(dist.assault, method = x)$ac
}
```

```
map_dbl(m, ac)
```

```
##   average   single  complete    ward
## 0.7379371 0.6276128 0.8531583 0.9346210
```


In order to get cluster identification, you will need to tell R where you want to cut your dendrogram.

```
h2=agnes(dist.assault,method="ward")
h2_clus <- cutree(h2, k = 2)
```

You can also compare different dendrograms....



```
dend1 <- as.dendrogram (h1.comp.eucl)
dend2 <- as.dendrogram (h2)
```

```
tanglegram(dend1, dend2)
```





You can also use the measures from the NbClust algorithm here, too....

```
NbClust(arrest.scal,distance="euclidean",method="complete",min.nc=2,max.nc = 4)
```

```
## *** : The Hubert index is a graphical method of determining the number of clusters.
##       In the plot of Hubert index, we seek a significant knee that indicates a
##       significant increase of the value of the measure i.e the significant increase in the
##       Hubert index second differences plot.
##
```



```
## *** : The D index is a graphical method of determining the number of clusters.
##           In the plot of D index, we seek a significant knee (the significant
##           second differences plot) that corresponds to a significant increase in
##           the measure.
```

```
##
```

```
## *****
```

```
## * Among all indices:
```

```
## * 9 proposed 2 as the best number of clusters
```

```
## * 9 proposed 3 as the best number of clusters
```

```
## * 5 proposed 4 as the best number of clusters
```

```
##
```

```
##           ***** Conclusion *****
```

```
##
```

```
## * According to the majority rule, the best number of clusters is 2
```

```
##
```

```
##
```

```
## *****
```

```
## $All.index
```

##		KL	CH	Hartigan	CCC	Scott	Marriot	TrCovW	TraceW	
## 2		26.3587	41.8949	11.5200	-0.3483	78.6988	725375.1	708.1600	104.6556	
## 3		0.0704	31.0737	18.9426	-2.0049	107.4372	918599.5	331.2367	84.3997	
## 4		393.5832	34.6264	5.4577	-0.7050	156.0842	617250.0	257.9307	60.1552	
##		Friedman	Rubin	Cindex	DB	Silhouette	Duda	Pseudot2	Beale	Ratkowsky
## 2		6.7099	1.8728	0.4031	1.0513	0.4048	0.5525	13.7698	1.8468	0.4404
## 3		8.1357	2.3223	0.3800	1.0433	0.3692	0.5918	20.0040	1.6098	0.4305
## 4		10.8610	3.2582	0.4334	1.0811	0.3160	1.0527	-0.3001	-0.1035	0.4155
##		Ball	Ptbiserial	Frey	McClain	Dunn	Hubert	SDindex	Dindex	SDbw
## 2		52.3278	0.6330	0.9663	0.5287	0.2637	0.0087	1.3745	1.3451	1.1226
## 3		28.1332	0.6315	0.8670	0.6873	0.2648	0.0092	1.2895	1.2025	0.7789

```
## 4 15.0388      0.5820 0.0448  1.3366 0.1622 0.0108  1.4085 1.0199 0.5029
##
## $All.CriticalValues
##   CritValue_Duda CritValue_PseudoT2 Fvalue_Beale
## 2           0.3773           28.0561           0.1300
## 3           0.4780           31.6758           0.1765
## 4           0.1265           41.4361           1.0000
##
## $Best.nc
##               KL           CH Hartigan           CCC Scott Marriot           TrCovW
## Number_clusters 4.0000 2.0000  4.0000 2.0000 4.000      3 3.0000
## Value_Index    393.5832 41.8949 13.4849 -0.3483 48.647 -494574 376.9233
##               TraceW Friedman Rubin Cindex           DB Silhouette           Duda
## Number_clusters 3.0000 4.0000 3.0000 3.00 3.0000      2.0000 2.0000
## Value_Index    -3.9886 2.7253 0.4865 0.38 1.0433      0.4048 0.5525
##               PseudoT2 Beale Ratkowsky           Ball PtBiserial Frey McClain
## Number_clusters 2.0000 2.0000 2.0000 3.0000      2.000 1 2.0000
## Value_Index    13.7698 1.8468 0.4404 24.1946      0.633 NA 0.5287
##               Dunn Hubert SDindex Dindex           SDbw
## Number_clusters 3.0000      0 3.0000      0 4.0000
## Value_Index    0.2648      0 1.2895      0 0.5029
##
## $Best.partition
##      Alabama      Alaska      Arizona      Arkansas      California
##      1           1           1           2           1
##      Colorado Connecticut Delaware      Florida      Georgia
##      1           2           2           1           1
##      Hawaii      Idaho      Illinois      Indiana      Iowa
##      2           2           1           2           2
##      Kansas      Kentucky Louisiana      Maine      Maryland
##      2           2           1           2           1
```

##	Massachusetts	Michigan	Minnesota	Mississippi	Missouri
##	2	1	2	1	2
##	Montana	Nebraska	Nevada	New Hampshire	New Jersey
##	2	2	1	2	2
##	New Mexico	New York	North Carolina	North Dakota	Ohio
##	1	1	1	2	2
##	Oklahoma	Oregon	Pennsylvania	Rhode Island	South Carolina
##	2	2	2	2	1
##	South Dakota	Tennessee	Texas	Utah	Vermont
##	2	1	1	2	2
##	Virginia	Washington	West Virginia	Wisconsin	Wyoming
##	2	2	2	2	2


4.2.1 Python for Hierarchical Clustering

To view the dendrogram in python, you can use the linkage command (caution!!! If you have a lot of data points, this is going to get really messy....in other words, I would not recommend it). This data set only has 50 observations, so it is not too bad to view. The method is the type of linkage and metric is the distance metric used.

```
from scipy.cluster.hierarchy import dendrogram, linkage

linkage_data = linkage(arrest_scal, method='ward', metric='euclidean')
dendrogram(linkage_data)

## {'icoord': [[5.0, 5.0, 15.0, 15.0], [25.0, 25.0, 35.0, 35.0], [10.0, 10.0, 30.0, 30.0], [15.0, 15.0, 35.0, 35.0], [5.0, 10.0, 20.0, 25.0], [10.0, 20.0, 30.0, 35.0], [15.0, 25.0, 35.0, 40.0], [5.0, 15.0, 25.0, 30.0], [10.0, 25.0, 35.0, 40.0], [15.0, 30.0, 40.0, 45.0], [5.0, 20.0, 30.0, 35.0], [10.0, 30.0, 40.0, 45.0], [15.0, 35.0, 45.0, 50.0], [5.0, 25.0, 35.0, 40.0], [10.0, 35.0, 45.0, 50.0], [15.0, 40.0, 50.0, 55.0], [5.0, 30.0, 40.0, 45.0], [10.0, 40.0, 50.0, 55.0], [15.0, 45.0, 55.0, 60.0], [5.0, 35.0, 45.0, 50.0], [10.0, 45.0, 55.0, 60.0], [15.0, 50.0, 60.0, 65.0], [5.0, 40.0, 50.0, 55.0], [10.0, 50.0, 60.0, 65.0], [15.0, 55.0, 65.0, 70.0], [5.0, 45.0, 55.0, 60.0], [10.0, 55.0, 65.0, 70.0], [15.0, 60.0, 70.0, 75.0], [5.0, 50.0, 60.0, 65.0], [10.0, 60.0, 70.0, 75.0], [15.0, 65.0, 75.0, 80.0], [5.0, 55.0, 65.0, 70.0], [10.0, 65.0, 75.0, 80.0], [15.0, 70.0, 80.0, 85.0], [5.0, 60.0, 70.0, 75.0], [10.0, 70.0, 80.0, 85.0], [15.0, 75.0, 85.0, 90.0], [5.0, 65.0, 75.0, 80.0], [10.0, 75.0, 85.0, 90.0], [15.0, 80.0, 90.0, 95.0], [5.0, 70.0, 80.0, 85.0], [10.0, 80.0, 90.0, 95.0], [15.0, 85.0, 95.0, 100.0], [5.0, 75.0, 85.0, 90.0], [10.0, 85.0, 95.0, 100.0], [15.0, 90.0, 100.0, 105.0], [5.0, 80.0, 90.0, 95.0], [10.0, 90.0, 100.0, 105.0], [15.0, 95.0, 105.0, 110.0], [5.0, 85.0, 95.0, 100.0], [10.0, 95.0, 105.0, 110.0], [15.0, 100.0, 110.0, 115.0], [5.0, 90.0, 100.0, 105.0], [10.0, 100.0, 110.0, 115.0], [15.0, 105.0, 115.0, 120.0], [5.0, 95.0, 105.0, 110.0], [10.0, 105.0, 115.0, 120.0], [15.0, 110.0, 120.0, 125.0], [5.0, 100.0, 110.0, 115.0], [10.0, 110.0, 120.0, 125.0], [15.0, 115.0, 125.0, 130.0], [5.0, 105.0, 115.0, 120.0], [10.0, 115.0, 125.0, 130.0], [15.0, 120.0, 130.0, 135.0], [5.0, 110.0, 120.0, 125.0], [10.0, 120.0, 130.0, 135.0], [15.0, 125.0, 135.0, 140.0], [5.0, 115.0, 125.0, 130.0], [10.0, 125.0, 135.0, 140.0], [15.0, 130.0, 140.0, 145.0], [5.0, 120.0, 130.0, 135.0], [10.0, 130.0, 140.0, 145.0], [15.0, 135.0, 145.0, 150.0], [5.0, 125.0, 135.0, 140.0], [10.0, 135.0, 145.0, 150.0], [15.0, 140.0, 150.0, 155.0], [5.0, 130.0, 140.0, 145.0], [10.0, 140.0, 150.0, 155.0], [15.0, 145.0, 155.0, 160.0], [5.0, 135.0, 145.0, 150.0], [10.0, 145.0, 155.0, 160.0], [15.0, 150.0, 160.0, 165.0], [5.0, 140.0, 150.0, 155.0], [10.0, 150.0, 160.0, 165.0], [15.0, 155.0, 165.0, 170.0], [5.0, 145.0, 155.0, 160.0], [10.0, 155.0, 165.0, 170.0], [15.0, 160.0, 170.0, 175.0], [5.0, 150.0, 160.0, 165.0], [10.0, 160.0, 170.0, 175.0], [15.0, 165.0, 175.0, 180.0], [5.0, 155.0, 165.0, 170.0], [10.0, 165.0, 175.0, 180.0], [15.0, 170.0, 180.0, 185.0], [5.0, 160.0, 170.0, 175.0], [10.0, 170.0, 180.0, 185.0], [15.0, 175.0, 185.0, 190.0], [5.0, 165.0, 175.0, 180.0], [10.0, 175.0, 185.0, 190.0], [15.0, 180.0, 190.0, 195.0], [5.0, 170.0, 180.0, 185.0], [10.0, 180.0, 190.0, 195.0], [15.0, 185.0, 195.0, 200.0], [5.0, 175.0, 185.0, 190.0], [10.0, 185.0, 195.0, 200.0], [15.0, 190.0, 200.0, 205.0], [5.0, 180.0, 190.0, 195.0], [10.0, 190.0, 200.0, 205.0], [15.0, 195.0, 205.0, 210.0], [5.0, 185.0, 195.0, 200.0], [10.0, 195.0, 205.0, 210.0], [15.0, 200.0, 210.0, 215.0], [5.0, 190.0, 200.0, 205.0], [10.0, 200.0, 210.0, 215.0], [15.0, 205.0, 215.0, 220.0], [5.0, 195.0, 205.0, 210.0], [10.0, 205.0, 215.0, 220.0], [15.0, 210.0, 220.0, 225.0], [5.0, 200.0, 210.0, 215.0], [10.0, 210.0, 220.0, 225.0], [15.0, 215.0, 225.0, 230.0], [5.0, 205.0, 215.0, 220.0], [10.0, 215.0, 225.0, 230.0], [15.0, 220.0, 230.0, 235.0], [5.0, 210.0, 220.0, 225.0], [10.0, 220.0, 230.0, 235.0], [15.0, 225.0, 235.0, 240.0], [5.0, 215.0, 225.0, 230.0], [10.0, 225.0, 235.0, 240.0], [15.0, 230.0, 240.0, 245.0], [5.0, 220.0, 230.0, 235.0], [10.0, 230.0, 240.0, 245.0], [15.0, 235.0, 245.0, 250.0], [5.0, 225.0, 235.0, 240.0], [10.0, 235.0, 245.0, 250.0], [15.0, 240.0, 250.0, 255.0], [5.0, 230.0, 240.0, 245.0], [10.0, 240.0, 250.0, 255.0], [15.0, 245.0, 255.0, 260.0], [5.0, 235.0, 245.0, 250.0], [10.0, 245.0, 255.0, 260.0], [15.0, 250.0, 260.0, 265.0], [5.0, 240.0, 250.0, 255.0], [10.0, 250.0, 260.0, 265.0], [15.0, 255.0, 265.0, 270.0], [5.0, 245.0, 255.0, 260.0], [10.0, 255.0, 265.0, 270.0], [15.0, 260.0, 270.0, 275.0], [5.0, 250.0, 260.0, 265.0], [10.0, 260.0, 270.0, 275.0], [15.0, 265.0, 275.0, 280.0], [5.0, 255.0, 265.0, 270.0], [10.0, 265.0, 275.0, 280.0], [15.0, 270.0, 280.0, 285.0], [5.0, 260.0, 270.0, 275.0], [10.0, 270.0, 280.0, 285.0], [15.0, 275.0, 285.0, 290.0], [5.0, 265
```



To do Agglomerative Clustering:

Note that you need to specify the number of clusters here. You are able to say None for the number of clusters, but then you need to specify the distance threshold (from dendrogram). Again illustrating that this is not the best method for large data sets.

```
from sklearn.cluster import AgglomerativeClustering
```


```
aggclus_py = AgglomerativeClustering(n_clusters = 2, affinity = 'euclidean', link  
arrest_hy = aggclus_py.fit_predict(arrest_scal)
```

I did not find a good package for divisive clustering in Python.

4.3 DBSCAN


DBSCAN is one of the most popular density-based clustering algorithm in use. Density Based Spatial Clustering Applications with Noise (DBSCAN) cluster points together that are “density reachable” (within a certain radius of each other). Users can specify this distance and the minimum number of points that will be used to create clusters (see <https://cran.r-project.org/web/packages/dbscan/vignettes/dbscan.pdf>)

```
scan1<-hdbscan(arrest.scal,minPts=4)
pca_ex=prcomp(arrest.scal,scale=F)
scan1data=cbind.data.frame(pca_ex$x[,1],pca_ex$x[,2],as.factor(scan1$cluster+1))
colnames(scan1data)=c("PCA1","PCA2","cluster")
ggplot(scan1data,aes(x=PCA1,y=PCA2,color=cluster))+geom_point()+ scale_fill_brewer
```




```
plot(scan1, show_flat=T)
```


HDBSCAN*



```
#res.dbscan=dbscan(arrest.scal,eps=1.2,minPts=4)
d=dist(arrest.scal,method = "canberra")
res.dbscan=dbscan(d,eps=1.2,minPts=4)
res.dbscan
```

```
## DBSCAN clustering for 50 objects.  
## Parameters: eps = 1.2, minPts = 4  
## Using canberra distances and borderpoints = TRUE  
## The clustering contains 2 cluster(s) and 27 noise points.  
##  
##    0    1    2  
## 27 10 13  
##  
## Available fields: cluster, eps, minPts, dist, borderPoints  
  
scan1data=cbind.data.frame(pca_ex$x[,1],pca_ex$x[,2],as.factor(res.dbscan$cluster))  
colnames(scan1data)=c("PCA1","PCA2","cluster")  
ggplot(scan1data,aes(x=PCA1,y=PCA2,color=cluster))+geom_point()+ scale_fill_brewer
```



###Nice visuals for PCA results:

<http://www.sthda.com/english/articles/31-principal-component-methods-in-r-practical-guide>

4.3.1 DBSCAN in Python

```
from sklearn.cluster import DBSCAN
from collections import Counter
```

```
db_py = DBSCAN(eps=1.2,min_samples=4).fit(arrest_scal)
```

```
db_py.labels_
```

```
## array([ 0, -1,  1,  1, -1, -1,  1,  1, -1,  0,  1,  1,  1,  1,  1,  1,  1,  1,
##         0,  1,  1,  1,  1,  1,  0,  1,  1,  1, -1,  1,  1,  1,  1,  0,  1,
##         1,  1,  1,  1,  1,  0,  1,  0,  1,  1,  1,  1,  1,  1,  1,  1],
##        dtype=int64)
```

```
set(db_py.labels_)
```

```
## {0, 1, -1}
```

```
Counter(db_py.labels_)
```


```
## Counter({1: 38, 0: 7, -1: 5})
```

4.4 Variable Clustering


```
telco=read.csv("Q:\\My Drive\\Data Mining\\Data\\TelcoChurn.csv",header=T)
telco[is.na(telco)] = 0
quant.var=telco[,c(3,6,19,20)]
qual.var=telco[,c(2,4,5,7:18)]

var.clust.h=hclustvar(quant.var, qual.var)
stab=stability(var.clust.h,B=50) ## This will take time!
```

Stability of the partitions



```
plot(stab)
```



```
h6=cutreevar(var.clust.h, 6)
```

4.4.1 Variable Clustering in Python

For variable clustering in Python, you can only have numeric values (unable to do this on categorical variables). HOWEVER, it does have a nice function involving ratio of R^2 .

```
from varclushi import VarClusHi
```

```
telco_py=pd.read_csv("Q:\\My Drive\\Data Mining\\Data\\TelcoChurn.csv")
```

```
telco_clus=telco_py[['SeniorCitizen','tenure','MonthlyCharges']]
```

```
telco_2=telco_py[['gender','Partner','Dependents','PhoneService','PaperlessBillin
```

```
telco_3 = pd.get_dummies(telco_2,drop_first=True)
```

```
telco_4=telco_3*1
```

```
telco_clus2=pd.concat([telco_clus,telco_4],axis=1)
```

```
telco_clus2.info()
```

```
## <class 'pandas.core.frame.DataFrame'>
## RangeIndex: 7043 entries, 0 to 7042
## Data columns (total 8 columns):
##  #   Column                Non-Null Count  Dtype
##  ---  ---
##  0   SeniorCitizen         7043 non-null   int64
##  1   tenure                 7043 non-null   int64
##  2   MonthlyCharges        7043 non-null   float64
##  3   gender_Male           7043 non-null   int32
##  4   Partner_Yes           7043 non-null   int32
##  5   Dependents_Yes        7043 non-null   int32
##  6   PhoneService_Yes      7043 non-null   int32
##  7   PaperlessBilling_Yes  7043 non-null   int32
## dtypes: float64(1), int32(5), int64(2)
## memory usage: 302.8 KB
```

```
varclus_py=VarClusHi(telco_clus2,maxeigval2=0.7,maxclus=None)
varclus_py.varclus()
```

```
## <varclushi.varclushi.VarClusHi object at 0x000001E830F3E6A0>
```

```
varclus_py.rsquare
```


##	Cluster	Variable	RS_Own	RS_NC	RS_Ratio
## 0	0	Partner_Yes	0.726338	0.144170	3.197620e-01
## 1	0	Dependents_Yes	0.726338	0.044599	2.864367e-01
## 2	1	MonthlyCharges	0.676075	0.061454	3.451351e-01
## 3	1	PaperlessBilling_Yes	0.676075	0.024502	3.320610e-01
## 4	2	gender_Male	1.000000	0.000256	2.221015e-16
## 5	3	PhoneService_Yes	1.000000	0.025753	2.279141e-16
## 6	4	SeniorCitizen	1.000000	0.052474	0.000000e+00
## 7	5	tenure	1.000000	0.100147	0.000000e+00